

Forecasting Constituents of the MSCI Minimum Volatility Index Through Logistic
Regression

A Thesis
Presented to
The Division of Statistics
Harvard College

In Partial Fulfillment
of the Requirements for the Degree
Bachelor of Arts in Statistics (Honors)

John A. Gilheany

November 6, 2017

Approved for the Division
(Statistics)

Professor Michael Parzen

David Kane

Acknowledgements

I want to thank Prof. Parzen and David Kane for all of their help.

Preface

This thesis explores a way of predicting index constituents using logistic regression.

Table of Contents

| | |
|--|-----------|
| Chapter 1: Introduction | 1 |
| 1.1 Background | 1 |
| 1.1.1 Exchange Traded Funds (ETFs) | 1 |
| 1.1.2 iShares MSCI Min Vol USA ETF | 2 |
| 1.1.3 Purpose | 2 |
| 1.1.4 Logistic Regression Model | 3 |
| 1.2 Literature Review | 3 |
| 1.2.1 Low-Risk Anomaly | 3 |
| 1.2.2 Low-Risk Anomaly Applications | 14 |
| Chapter 2: Mathematics and Science | 21 |
| 2.1 Math | 21 |
| 2.2 Chemistry 101: Symbols | 21 |
| 2.2.1 Typesetting reactions | 22 |
| 2.2.2 Other examples of reactions | 22 |
| 2.3 Physics | 22 |
| 2.4 Biology | 22 |
| Chapter 3: Tables, Graphics, References, and Labels | 23 |
| 3.1 Tables | 23 |
| 3.2 Figures | 24 |
| 3.3 Footnotes and Endnotes | 26 |
| 3.4 Bibliographies | 26 |
| 3.5 Anything else? | 28 |
| Conclusion | 29 |
| Appendix A: The First Appendix | 31 |
| Appendix B: The Second Appendix, for Fun | 33 |
| References | 35 |

List of Tables

3.1 Correlation of Inheritance Factors for Parents and Child 23

List of Figures

| | | |
|-----|--|----|
| 1.1 | Capital Asset Pricing Model | 9 |
| 1.2 | Information Ratio | 11 |
| 3.1 | Reed logo | 24 |
| 3.2 | Mean Delays by Airline | 25 |
| 3.3 | Subdiv. graph | 26 |
| 3.4 | A Larger Figure, Flipped Upside Down | 26 |

Abstract

The low-risk anomaly has created opportunities for arbitrage in the financial markets. As Baker et al. discuss in “Benchmarks as Limits to Arbitrage: Understanding the Low-Volatility Anomaly,” low-volatility and low-beta portfolios outperform and high-volatility and high-beta portfolios by a factor of several times due to benchmarking and lottery-preferences. The iShares MSCI USA Minimum Volatility (USMV) is an ETF tracking a minimum volatility index that was used to find data and will be used for trading arbitrage. Frazzini et al. discuss arbitrage opportunities by quantitative focused funds like AQR in “Betting Against Beta”, and this thesis explores a more advanced type of index front-running as a potential arbitrage opportunity. Data was collected from USMV from its inception in October 2011, and from EUSA, the parent ETF of USMV, from the same period until December 2016. 52-week trailing beta, 52-week trailing volatility, lagged price/book, and current index membership were calculated, and a regression model was run to quantify the relationship between current index membership and these four variables. In the model, a probabilities of index membership were calculated and an optimal cutoff was calculated to which the model would be 95% accurate of its findings of a stock to be in or out of USMV, given the historical data. Backtesting with prior data showed with a model accuracy of 95%, arbitrage opportunities of X% could be collected after each rebalancing.

Chapter 1

Introduction

1.1 Background

The iShares Edge MSCI USA Minimum Volatility (USMV) Exchange Traded Fund (ETF) is designed to track the investment results of the MSCI Minimum Volatility USA index, which is composed of stocks with a lower volatility than the general market. This can provide investors with exposure to a portfolio with less risk than many alternatives, and historically has declined less in value than the broader market during economic downturns. The ETF is comprised of 189 holdings, and is rebalanced two times per year, with the intention of mirroring the changes made by the index. The purpose of this dissertation is to create a logistic regression model that can accurately predict which stocks will be added or removed from this ETF before rebalancing occurs, and understand what factors are involved. The model will take into account volatility attributes of each stock, as well as others potentially significant predictor variables from prior studies. An accurate model will allow for arbitrage investment opportunities.

1.1.1 Exchange Traded Funds (ETFs)

An Exchange Traded Fund (ETF) is a collection of stocks and/or bonds in a single portfolio, that is traded on a major exchange just like a stock is (Hayes, 2017). As a result, the price of an ETF fluctuates on a regular basis. Exchange Traded Funds generally have more liquidity and less fees when compared to other alternatives instruments like mutual funds. Owning an ETF can allow investors to minimize risk, since owning an ETF is comparable to owning a little bit of many different stocks. This diversification comes at lower costs and less effort for investors as well.

ETFs can also track an index, commodity, bond, or basket of all of the above. Unlike an ETF, which is publicly traded, an index is not. The goal of the USMV ETF is to track the MSCI Minimum Volatility USA index, and this is more complicated than it seems. In addition to tracking this index, the ETF aims to mirror returns of the index, and any difference is called tracking error. The tracking error is often very small, and can be around a tenth of a percent. This error can come from indices being market capitalization weighted, meaning that price fluctuations of each stock lead to

the weighting being changed by a ratio of its market cap against the market cap of all stocks in the index (Fontinelle, 2009). With these stocks weightings in the index constantly changing and people buying in and out of ETFs constantly, it is hard to track performance entirely accurately. However, ETFs very closely follow indices, as their tracking errors are generally quite small. Thus, although ETF data is not the same as index data, they are extremely comparable.

1.1.2 iShares MSCI Min Vol USA ETF

The iShares Edge MSCI Min Vol USA ETF (USMV) is a Blackrock-managed ETF that tracks the investment results of the MSCI Minimum Volatility USA index. The MSCI Minimum Volatility USA index constituents come from the MSCI USA Index, which are roughly comprised of the top 600 US stocks by market cap. This minimum volatility index is intended to have a lower beta, lower volatility, lower cap bias, and contain more stocks with less risk than its parent index, which contains US mid-cap and large-cap stocks. The index is rebalanced twice a year, on the last trading days of May and November. The index typically has around 180 constituents, with an average of 20 new additions and 14 deletions every 6 months when rebalancing occurs. Over the last five, years, the number of additions has ranged from 12 to 25, while the deletions have been between 10 and 19. Changes to the index are usually announced nine trading days before they are set to take place (BlackRock, 2017).

Using the Barra Open Optimizer, USMV creates a minimum variance portfolio of low risk stocks, as a subset from its parent index of USA large-cap and mid-cap stock. Using this estimated security covariance matrix, the MSCI Minimum Volatility Index is the product of the lowest absolute volatility, considering the constraints (MSCI, 2013). Moreover, these additions are simply a relabeling of existing stocks in the parent index, and do not include new additions to the parent index. The low-risk stocks chosen to be in USMV are determined by a set of constraints, like maintaining a certain sector or country weight relative to the parent index.

There are many specific constraints to this index. The first is that an individual stock cannot exceed 1.5% or 20 times the weight of the stock in the parent index. The minimum weight of a security in the index is also capped at 0.05%. USMV also aims to keep the weight of specific countries within a 5% range of the weight in the parent index, or 3 times the weight of the country in the parent index. Sector weights of USMV cannot deviate more than 5% from the sector weights in the parent index. One way turnover of the index is also maxed at 10%. Thus, taking into account these constraints, the Barra Open Optimizer creates the lowest absolute volatility portfolio possible (Wynne, 2016).

1.1.3 Purpose

As mentioned, the purpose of this thesis is to create a model to that will predict rebalancing of stocks in the Min Vol index, and consequently the USMV ETF, before it actually happens. There is substantial price movement whenever a stock is added or removed from a large ETF, like USMV. When a stock is added to the index, the ETF

will buy large amounts of that stock, increasing the demand, and consequently, the market price for that stock. If the stock is bought in advance of this large purchase, then the investor can enjoy pretty immediate price appreciation in the stock. Moreover, if a stock is removed from the Min Vol index, the USMV ETF will sell all current holdings of the stock, which will increase the supply of the stock, driving down market price of the stock. If one were to short this stock before that happened, he/she can also profit from that event.

A phenomena known as ETF front-running is similar to what this paper hopes to accomplish, but is one step behind. ETF front-running occurs when traders buy or sell stocks in advance of ETF managers after they announce an exit or entrance of a position (MacDonald, 2009). There is typically a slight lag between an announcement of an ETF to add or remove a position, and the actual purchase or sale of this position. By acting quickly, traders can scalp profit by buying a stock before an ETF does, and selling it to them later at a slight profit, or short-selling a stock before an ETF exits the position, and then buying it back at the lower price. The thesis will take this one step farther, and try predict the stock addition or deletion before announcement. This will allow traders to similarly front-run the index, but they will do so before the market is able to react, leading to larger profit opportunities.

1.1.4 Logistic Regression Model

These goals of this paper will be achieved by creating a logistic regression model, which will be transformed to calculate a probability of a stock being in the out of the index. A logistic regression model will be used since the dependent variable is dichotomous. The predictor variables will include 52-week trailing volatility, 52-week trailing beta, price/book ratio, and whether or not the stock was in the index 6 months before, during the previous rebalancing. These attributes were chosen after looking at the historical literature and understanding of the minimum volatility index.

1.2 Literature Review

1.2.1 Low-Risk Anomaly

The Low-Risk Anomaly stems from the observation that portfolios of low-risk stocks have higher risk-adjusted returns than portfolios of high-risk stocks. This challenges a widely regarded financial principle that investing in higher-risk stocks will generally result in higher expected return. In several studies, portfolios of high-risk stocks and low-risk stocks were constructed and rebalanced regularly to reflect these characteristics, and the low-risk portfolios outperformed the high-risk portfolios by a factor of several times, over long-periods of time including 1929-2015 (Vliet & Koning, 2016) and 1968-2008 (Baker, Bradley, & Wurgler, 2011).

One the explanations is the need for money managers to be compared to a benchmark, such as an index like the S&P500. This is reasonable, as many fund managers are charging fees to manage money, and need a way to prove themselves and

their abilities. By outperforming an index, a fund manager is able to generate alpha, and presumably raise money money or charge higher fees. By underperforming an index, the fund managers has a hard time justifying fee charged to clients, since they could just invest in the index passively for little to no fee. Thus, much of the risk for them is relative, coming from potentially underperforming a benchmark. Moreover, with this doubling of assets under active management from 30% to 60% in 1968-2008, the low-risk anomaly intensified. Another metric of fund manager performance is through the “information ratio” (IR), or the expected return difference between the manager and the expected return of the S&P 500, divided by the volatility of this return difference. The goal of an investment manager is to maximize this number, best as possible, through picking stocks that will outperform the market. These help to create a greater demand for higher-risk stocks, while discouraging investments in low-beta, low-volatility stocks, and ultimately increases the market’s appetite for risky stocks with high reward potential, driving up price and driving down expected return.

In addition to the focus on relative rather than absolute risk, money managers also often focus on single period returns, often as short as a month, which aim to remove the effects of compounding. This helps differentiate each manager’s stock picking abilities, but ignores the real-world and significance of compounding. Humans also have a predisposition for the lottery affect, which is increased interest in stocks with high skew - that is high upside potential. Risk can also be decomposed into macro and micro-effects - that is, looking for and comparing the risk-return characteristics of stocks with different risk profiles, but similar country and industry risks. It turned out that the micro-effects, those that were stock-specific, were statistically significant at generating alpha, while the macro-effects, those that were country and industry-specific was not. These literature reviews also examine both volatility and beta as a measure of risk, and suggest that beta is not an adequate risk metric. In fact, though beta and volatility are obviously very correlated, beta appears to be more related to this low-risk anomaly than volatility. This can have significant implications on the significance of the predictor variables in the logistic regression model in this paper. In the following section, these findings will be highlighted and discussed in more detail.

High Returns from Low Risk, By Pim van Vliet and Jan De Koning

Assuming an efficient market, one of the most widely accepted tenants of investing is that one will be receive higher reward for taking on more risk. Presumably, if this were not true, nobody would partake in higher risk investments. Some risks to consider when making an investment include those with respect to the market, liquidity, credit, inflation, and FX (Ontario, 2017). When quantifying risk for a company, one can look at the standard deviation of the stock price. By looking at the historical dispersion of data from the mean, or normal returns, one find the stock’s volatility. These calculations are based only on the price fluctuations of the stock, and no other external factors (Investopedia, 2015). Volatility is also one of the best indicators of bankruptcy. The other common form of risk measurement in finance is beta, which measures the stock’s price volatility with regard to the market. The market is typically a benchmarked index relevant to the stock; for a large-cap US

stock, the associated index could be the S&P 500. Beta is calculated by taking the covariance of the stock returns and the market returns, then dividing by the variance of the market. This yields a coefficient, which can be interpreted: a beta value of 1 indicates the stock price and market move together identically, a beta value of less than 1 indicates the stock is less volatile than the market, and a beta value of greater than 1 indicates the stock is more volatile than the market. Beta values can be negative as well, and hold the same interpretation as positive beta values, with the difference being that the stock price and market move in opposite directions. For example, a beta of -1 would mean the stock and market have the same volatility and changes in price, but that the performance is inversely related. Thus, the first sentence of this paragraph can have many meanings, but can be interpreted as saying that stocks with higher beta or volatility will have higher expected return.

In their book, Jan de Koning and Pim Van Vliet set out to investigate the question: Do high-volatility stocks return more than low-volatility stocks? The authors constructed high-volatility and low volatility portfolios, and compared the two over a 86-year time span, from 1929-2015. Over this period, low volatility stocks outperformed high volatility stocks by a factor of 18, excluding inflation and transaction costs. If both portfolios started off with \$100 in 1929, the low-volatility portfolio end value would be worth \$395,000, while the high-volatility portfolio would be worth just \$21,000. In reality these values would both be much smaller if the costs of trading the stocks and inflation were considered, but excluding them is reasonable as the effect on both portfolios would be pretty similar. The low-volatility portfolio returned 10.2% annually whereas the high-volatility stocks returned just 6.4% annually. This annual difference of 3.8% is striking, and when considering compounding over an 86-year period, explains why the low-volatility portfolio's ending value was over 18 times that of the high-volatility portfolio. In this study, the annualized volatility of the low volatility portfolio was 13%, and the annualized volatility of the high volatility portfolio was around 2.5 times that, at 36% (Vliet & Koning, 2016).

These findings present an anomaly in the field of finance, and beg the question of why anybody would invest in high-volatility stocks in the first place. Moreover, one may wonder how it is possible that a portfolio of low-volatility stocks can outperform high-volatility stocks over a long period of time. This anomaly has several explanations.

It is important to understand that a higher volatile stock or portfolio will move in greater magnitude than the underlying market. This holds true for both downside and upside scenarios. When the market increases in value, a high volatility stock will increase in excess of this. For example, if the market increases by 20% in one year, a high volatility portfolio would reasonably appreciate by more than 20% over the same period. However, when the market declines, a high volatility portfolio will decrease in excess of this amount. For instance, if the market decreases by 20% in one year, a high volatility portfolio would reasonably depreciate by more than 20% over the same period. Lower volatility portfolios would react in similar fashion, just with a lesser magnitude with respect to the market. With this in mind, one way the low volatility portfolio is able to outperform the higher volatility portfolio is by losing less during times of financial stress. The authors conveniently started their analysis right at the beginning of the most severe economic depression in American

history, the Great Depression. The Great Depression began in 1929 and eroded away around 80% of the market's value by the time the recovery began in 1932 (History, 2010). With this in mind, by 1932 the high volatility portfolio declined by over 80%, while the low volatility portfolio declined by less than 80%. More specifically, the high volatility portfolio shrunk from \$100 in value to just \$5, while the low volatility portfolio shrunk from \$100 in value to \$30. Thus, the results of this paper can be taken with a grain of salt, since although the portfolios each started off with the same amount of money in 1929, the low volatility portfolio was worth 6 times as much as the high volatility portfolio just four years later. However, this is an expected consequence of the high volatility portfolio, so these results should not be discredited. With this being said, since the low volatility portfolio was able to lose less money during market downturns, it is able to grow capital more effectively than the high volatility portfolio. To illustrate this, the portfolio values can be considered in 1932. As Benjamin Graham, a famous value investor, once noted the mathematical fact that "once you lose 95% of your money, you have to gain 1,900% just to get back where you started" in his classic book "The Intelligent Investor" (Graham, 1965). Likewise, when the high volatility portfolio lost six times as much value as the low volatility portfolio, it would have to outperform the low volatility portfolio by significantly more than 600% in order to return to the same value.

Thus, it seems very counterintuitive that fund managers and investors do not only invest in low volatility stocks, as it appears that these stocks will outperform high volatility stocks in the long run. Part of understanding why this is not a commonplace investment strategy for many comes from interpreting what risk is defined as by the financial community. Risk is not necessarily analogous to volatility to them, or even to losing money. For many fund managers, risk comes from underperforming a benchmark. David Blitz, the head of quantitative equity research at Robeco, discussed the need for benchmarking the performance of investment managers. Many managers, especially of actively managed funds, command handsome compensation in exchange for their investment acumen and diligence. Many hedge funds have a "Two and Twenty" compensation structure, where the managers charge a 2% fee on total assets under management and take an additional 20% of profits (Investopedia, 2017). Given the large amount of fees clients are paying, it is reasonable to believe they expect to receive a greater return than if they had invested in a passive, market tracking index. These investment professionals need to prove to their bosses and clients that they are above average in their job. For example, if a hedge fund manager returns 10% in one year, but the market returns 10% that same year, the client will be upset, as they are now receiving a smaller return, given the hefty fees. In this case a \$100 investment in a market tracking index, would return approximately \$10 pre-tax, given there are little to no management fees. However, \$100 invested in a hedge fund with a "Two and Twenty" structure would yield the same initial \$10 pre-tax return, but the client would pay 2% of \$100 (\$2) in management fees, and 20% of the \$10 profits (\$2) in additional compensation. This would leave the client with a total of \$6 return, or 6%, after all of the fees are paid out. As a result, given the fee structure and client demands, active money managers are often compared to a market benchmark like the S&P 500. More importantly, they are expected to outperform these benchmarks

substantially. Benchmarking also helps add perspective to a manager's returns. If a manager returns 20% in one year when the market returns 10%, he/she will have many happy clients, but if the manager returns 20% in the same year that the market returns 30%, clients will not be as pleased. This is a concept known as "relative" risk.

Many investors, like individuals saving for retirement, primarily care about absolute risk, which is the total amount of money that is gained or lost due to overall stock movements, with regard to the starting amount of money invested. They will check their portfolio's total performance without much concern for the exact market return. If their portfolio gains 20% in one year, they will be content, even if the market increases 30% in the same period. For these investors, the horizon is long-term so the short-term performance isn't as important for them as it is for fund managers who may be trying to raise more capital or justify high fees from clients. Volatility, in itself, captures these changes in the price of a stock, and is an absolute risk measurement. Many institutional investors do not look at risk on an absolute level, as a retiree or mom and pop investor may, but instead look at the risk of a portfolio with respect to market or some other widely accepted benchmark. For these investors, the risk is not so much about losing money, rather is more centered around lagging the market or their peers. Investing is very much a relative game. To further illustrate this idea, for a fund manager if a portfolio drops 20% while the market drops 40%, this is seen as a much better outcome than if a portfolio goes up 20% while the market goes up 40%. In the former, the manager lost money, but outperformed the market by 20%. In the latter, the manager made money, but lagged the market by 20%. This concept can be hard to fully grasp due to the natural bias towards focus on absolute risk. Several other papers have tried to explain this seemingly misunderstood phenomena, including Karceski in 2002, where it was noted that an extrapolation bias could cause mutual fund managers to care more about overperforming in a bull market, than underperforming in a bear market (Karceski, 2002).

Another explanation of the low-volatility anomaly is an increased focus on returns over short-term periods by many researchers and investors. By focusing on "single period returns", which in most academic studies is just a one-month period, the significance of compounding is removed. This is more of an arithmetic way to calculate returns, where each month's return can be averaged, for example. Mathematically, this is not the correct way to calculate a return since it does remove the affect of compounding, but is common in a way to compare fund managers. Moreover, when done over very short time periods (like a year), the effect of compounding is not as significant as it is for very long-term periods. To illustrate this, the following scenario can be considered: in one month a portfolio worth \$100 drops 50% to \$50, then the next month increases 50% to \$75. The investment return is dependent on how one divides the time period. Looking at it on a monthly basis, even though the portfolio lost \$25 in value, the net return would be $-50\% + 50\%$, or 0% (with focus on single period returns, not accounting for compounding). However, looking at it on a longer term basis, the net return was -25% , as the \$100 portfolio ended up losing \$25 in value. By not fully including the magic "return upon return" effect of compounding, the high-volatility portfolio in this book performs more than 6% better per year.

In addition to the reasons mentioned, there are several psychological reasons why

some investors are not attracted to low-risk stocks. Eric Falkenstein, a renowned author in the low volatility investing realm, wrote that “envy is at the root of the investment paradox.” Some investors simply don’t recognize the significance of compounding returns. Many others do, but are unable to utilize the paradox due to relative risk and career pressures. Analysts who choose big winners are more likely to get recognized and promoted than those who pick safer stocks with lower upside potential; funds that pick the right high risk stocks that turn out to be major home-runs will see more reward as an increase in AUM, and consequently an uptick in management fees. Moreover, some people do not invest in low risk stocks because they have less of an appeal than high risk stocks, where investors think they can make a lot of money easily and quickly. Even the news will have a bias towards reporting about and covering stocks that have become big winners. One famous big winner is Amazon - a \$5,000 investment in 1997 would be worth \$2,400,000 today, or an increase 49,000% (Berger, 2017). With all the excitement and reporting to this day on big winners like Amazon, many people forget about the number of similar technology companies that failed and are worth nothing today. These high-risk stocks are more “sexy” and have a “lottery ticket” element that attracts investors with the appeal of a big payday.

Benchmarks as Limits to Arbitrage: Understanding the Low-Volatility Anomaly, by Malcolm Baker, Brendan Bradley, and Jeffrey Wurgler

In the paper, the authors performed a study similar to that done by van Vliet and De Koning, except using 41 years of CRSP data, ranging from January 1968 – December 2008. Once again, it is important to note the ranges of dates used. The Great Recession began in 2007 after the bursting of the subprime mortgage bubble, leading to the collapse of several large, investment banks, and the government bailout of many others. This caused market indices like the Dow to drop from a high in 2007 of 14,164.43 to 8,776.39 by December 2008, representing a decrease in over 38% during that period (Amadeo, 2017). Thus just as van Vliet and De Koning started their sample right before the Great Depression to help amplify their results, Baker, Bradley, and Wurgler ended their sample after the Great Recession. Though this may help amplify the results, this does not take away from the significance of the findings.

Low-volatility, high-volatility, low-beta, and high-beta portfolios were constructed using the top 1,000 stocks by market capitalization, and then calculating each stock’s five-year trailing volatility or beta. A dollar investment in the low-volatility portfolio in 1968 appreciated to \$59.55 by 2008, or \$10.12 in real terms when accounting for inflation. On the other hand, the highest-volatility portfolio went from a dollar in value to 58 cents from 1968-2008, with a real value of just around 10 cents when considering inflation. Thus the low-volatility portfolio outperformed the high-volatility portfolio by over 100 times both in terms of nominal and real value. When using beta as a measure for risk, the finding was very similar. In the lowest-beta portfolio, a dollar grew to \$60.46 in nominal value throughout the 41-year period, or \$10.28 in real value after considering inflation. The highest-beta portfolio grew from a dollar to \$3.77 in nominal terms, or \$0.64 in real terms after accounting for the effects of inflation. Thus, the low-beta portfolio outperformed the low-beta portfolio by a factor

of 16, in nominal and real terms. One interesting observation here is though both the low-beta and low-volatility portfolios outperformed the high-beta and high-volatility portfolios, respectively, the discrepancy was far more pronounced in the case of the the volatility portfolios. Another significant observation was that investors who owned either the high-beta or high-volatility portfolio in 1968, would have lost money, when accounting for the effects of inflation by 2008. Another important thing to note is that this was just the case for large-cap companies, as these portfolios were constructed from the top 1,000 stocks in terms of market cap. The fact that this anomaly was observed for large-cap companies is quite impressive, given there are generally a lesser degree of mispricings in that realm than with small-cap companies, because many small-cap stocks are not big enough for institutional investors to purchase. Thus, this anomaly would be larger if done for small-cap stocks instead. In addition, the portfolio end values assumed no transaction costs, which in reality should be considered. The high-beta and high-volatility portfolios cost more to rebalance on a monthly level, as was done in the paper, than the low-beta and low-volatility portfolios, indicating this anomaly is actually more pronounced than initially reported. Baker et al. noted that while high-beta portfolios outperformed the low-beta portfolios in up markets and underperformed the low-risk portfolios in down markets, that the low-beta anomaly persisted in both situations. On a market-adjusted basis, low-beta consistently generates high alpha. This was consistent with prior research (Pettengill, Sundaram, & Mathur, 1995).

These findings are not new or revolutionary, and have been observed in many previous academic articles. Black, Jensen, and Scholes evaluated some of the assumptions and effectiveness of the Capital Asset Pricing Model (CAPM) in 1972 (Jensen, Black, & Scholes, 1972). The CAPM is a model that quantifies the relationship between risk and expected return for a stock, and considers investors should be compensated for the time value of their money and risk they are taking. This is calculated as the risk free rate plus the beta times the difference between the risk free rate and expected market return (Brennan, 1989):

$$\bar{r}_a = r_f + \beta_a(\bar{r}_m - r_f)$$

Where:

r_f = Risk free rate

β_a = Beta of the security

\bar{r}_m = Expected market return

Figure 1.1: Capital Asset Pricing Model

This will give the expected return of the stock, and anything in excess of this will be considered alpha. Black, Jensen, and Scholes were able to find that expected excess return, alpha, but not strictly proportional to Beta, as it is in the CAPM. In 1975, Haugen and Heins similarly found that there is little support for the idea that risk premiums have manifested themselves in realized rates of return (Haugen & Heins,

1975). In fact, they point out that the relationship between risk and return is much flatter than it is in the CAPM. In 1992 Fama and French famously declared that beta was dead after finding a flat relationship between beta and return (Fama & French, 1992). Many additional papers and research has added evidence for disproving the Capital Asset Pricing Model (CAPM), and even suggest that beta may not be the correct measure of risk, and that the relationship between risk and return is not what many people believe it to be (Mullins, 1982). However, other models relating risk and return, have had difficulty gaining acceptance and widespread usage in the finance industry. This paper is somewhat unique, in that Baker et al. does not try to prove what has already been shown in several academic works of literature, but instead offer several explanations for why this discrepancy exists.

One theory that is explored in detail is an investor's irrational preference for high-volatility stocks. Many investors have a natural preference for lotteries, even though there is a general aversion towards loss. If a stock has a positive skew, which is defined as a larger probability of a large positive payoff than probability of a small payoff, investors typically are very interested. Though skew is not the same as volatility, through their research, Boyer, Mitton, and Vorkink make a strong case for how expected skewness is a proxy for volatility through their findings that expected skewness assists in explaining the observation that stocks with high idiosyncratic volatility have low expected returns (Boyer, Mitton, & Vorkink, 2009). Another idea is representativeness, or that Bayes' rule and probability theory are often not natural to people, even the most seasoned investment professionals. One example of this is selectively looking at a few speculative investments that have turned out to be massive successes, without considering the numerous failures. As mentioned earlier, the news will focus on Amazon's great success over the past twenty years, but will not focus as much on all of the tech stocks that became worthless after the bubble burst. By not separating and considering all winners and losers, the average investor may be inclined to overpay for a riskier stock. Overconfidence has also been tied to a preference for volatile stocks; optimists are generally more aggressive than pessimists. In a study, people were asked questions and how certain they were of their responses and it appeared many did not have an understanding of probability (Fischhoff, Slovic, & Lichtenstein, 1977). Estimating the heat of a candle flame is very difficult, so for someone to say they are 80% sure of their answer is impressive, yet hard to believe. This can apply when people are asked for a confidence interval of the population of a city. Many times the person would be confident and give a very narrow range of people. This same concept applies when valuing stocks and evaluating certain investment opportunities.

Moreover, as mentioned previously, the need for benchmarking, especially among institutional investors, is also believed to heavily play into the anomaly. In fact, in the first year of this study, institutional investors managed 30% of all money, but by 2008, the final year of the study, this figure increased to 60%. With this doubling of assets under active management, the anomaly being discussed intensified. This still begs the question as to why institutional investors do not buy more low-volatility stocks, and the answer has to do with benchmarking. One explanation involves how the typical fund manager is measured for their performance. One common metric

is the “information ratio” (IR) which is the expected return difference between the manager and the expected return of the S&P 500, divided by the volatility of this return difference (tracking error) (Goodwin, 1998):

$$IR_P = \frac{\bar{R}_P - \bar{R}_B}{\hat{\sigma}_{P-B}}$$

\bar{R}_P represents the return of the portfolio for the time period under measurement

\bar{R}_B represents the return of the benchmark

$\hat{\sigma}_{P-B}$ is the standard deviation of the difference in returns between the portfolio and its benchmark

Figure 1.2: Information Ratio

The goal of an investment manager is to maximize this number, best as possible, through picking stocks. In fact, over 61% of U.S. mutual fund managers are benchmarked against the S&P 500, while over 94% are benchmarked to some U.S. index benchmark (Sensoy, 2009). Moreover, SEC rules require mutual funds to compare their performance to some benchmark (<https://www.sec.gov/reportspubs/investorpublications/investorpubsinwsmfhtm.html>). This intuitively makes sense, as it allows investors to assess the skill and ability of managers in an unbiased way, and also allows fund managers a chance to differentiate themselves. However, this makes institutional investors, who are managing the majority of the money in the US, less likely to buy low-volatility stocks, leading to higher prices and lower expected returns for the high-volatility stocks and further exacerbating the low-risk anomaly.

Investment managers without leverage will try to find mispriced stocks with a beta very close to market risk (beta of 1), overweighting positive-alpha stocks while underweighting negative-alpha stocks. When comparing the Sharpe ratio of large cap stocks for a low-volatility portfolio, it was found to be quite high at 0.38. However, IR was a very low 0.08, showing this would be very tough for a fund manager to invest in. To provide a comparison, during 1968-2008, the top value strategy portfolios had an IR of 0.51, and top momentum strategy portfolios had an IR of 0.64. This is extremely high compared to the IR of low-volatility stocks in this period, which ranged from 0.08 to 0.17, showing these constructed low-beta portfolios would be used by any fund manager. While Beta and volatility are undoubtedly very correlated, this paper showed that beta is more related to the anomaly than volatility, especially with large cap stocks, which is what most fund managers disproportionately focus their investments in.

Overall, irrational investor preference for lotteries and high volatility stocks, as well as investment managers' focus on benchmarks and IR, flatten and eventually invert the relationship between risk and return in the long-run. Moreover, it has been shown by Baker et al., and prior observations that the anomaly intensified

with the increase in assets under active management of fund managers in the U.S. These reasons together have led to the findings that low-beta and low-volatility stocks have outperformed high-beta, high-volatility stocks from 1968-2008, in part due to combining great returns with low downturns. Investor preference for “lotteries” and a bias of overconfidence creates a higher demand for higher-volatility securities, and the need to benchmark creates a greater demand for higher risk stocks, while discouraging investments in low-beta, low-volatility stocks. This understandably, increases the market’s appetite for risky stocks with high reward potential, driving up price and driving down expected return. These reasons appear perennial, so the anomaly is unlikely to cease to exist in the near future.

The Low-Risk Anomaly: A Decomposition into Micro and Macro Effects

The low-beta anomaly, which as discussed, relates to the outperformance of low-beta stocks, can be broken up into micro and macro effects. The micro effects include picking low-beta stocks, while keeping country and industry risk the same, while the macro effects involve picking low-beta countries or industries, while keeping stock-specific risk the same. In this paper, the micro effects were recorded by creating portfolios of equity longs and shorts, holding forecasted country and industry risk constant. The macro effects were observed by constructing long-short portfolios of various countries and industries, holding forecasted stock-specific risk constant. Studying a number of stocks within 29 industries and 31 different developed countries, the macro and micro effects were observed, and both together were shown to play an important role in the low risk anomaly.

Looking at 31 developed countries including Canada, France, Germany, Japan, and Singapore, the paper worked to decompose the low risk anomaly into country and stock specific effects. Similar to the industry findings, country-beta was able to predict stock betas to a certain extent, but not as well as historical stock betas were. Looking only at country betas yielded around half the risk reduction and two-thirds the risk adjusted return improvement, as compared to stock betas. This study implies that predicting risk of individual stocks is in itself very hard when only given data on country or industry risk, but when given all the data can have much more predictive power.

It was found that using industry beta to predict future stock betas was possible, but not as effective as just using historical stock betas. Industry beta information without stock information does improve risk-adjusted returns, just not to the same extent as it does with stock information. The paper also goes into detail trying to isolate pure industry effects and pure stock effects. Pure industry effects are the average differences between high and low beta industries, while holding constant stock risk. Pure stock risk is the opposite of that, calculating the average difference between high beta and low beta stocks, keeping industry risk constant. In the end, finding low-risk portfolios using selection of low risk stocks keeping industry constant was around four times more effective than using industries and keeping stock risk constant. Using the historical betas of both together, however, has more predictive power than either one alone.

Micro-selection of stocks, holding country and industry risks constant, was shown to be able to significantly reduce risk without a significant decrease in return. In some cases, high-risk stocks within particular industries were able to be distinctly identified, and they typically had similar returns when compared to low-risk stocks in the same industry and country. Macro-selection, holding stock-specific risks constant, was shown to lead to increases in return with small differences in risk especially with regards to the country chosen. High-risk countries were found to have distinctly lower returns than low-risk countries. These findings hold significant investment opportunity, due to the implication that people seeking arbitrage opportunities through mispricing of macro-effects like industry and sector, or through ETFs might not be as successful as exploiting the risk reduction opportunities stemming from micro-effects. While both the micro and macro-effects led to higher CAPM alpha by reducing risk and increasing returns, only the micro effects were found to be significant, whereas the macro-effects of country selection and industry selection were found to not be statistically significant. There is more of an arbitrage opportunity exploiting the micro-effects of individual stock selection than the macro effects. Even though the macro-effects have many limitations in practice, micro-effects are also limited due to leverage restrictions and benchmark mandates.

Understanding Defensive Equity, by Robert Novy-Marx

Defensive equity strategies generally will aim at including more safe/defensive stocks than risky/aggressive stocks, often with respect to a stock's volatility or beta. This strategy has been becoming very popular as of late, due in part to equity markets that have suffered two recent, severe downturns, negative nominal returns in the first decade of 2000, and literature proving a weak or negative relationship between risk and return. Lots of prior literature has looked at the relationship of performance relative to size and value, like some of the papers discussed above, but many have not done a deeper dive into the effect of specific factors, such as profitability. In fact, it has been shown that high profitability is the most significant predictor of low volatility, inherently causing these defensive equity strategies to indirectly tilt towards profitability, in addition to more known factors like size and value. Size is very important, as small growth stocks are typically underweighted in these strategies, while large value stocks are overweighted. Likewise, value stocks are typically overweighted, while growth stocks are underweighted. Thus, the performance of defensive equity strategies can be explained through accounting for size, value, and profitability.

In terms of performance, defensive equity strategies, which are defined as low-volatility or low-beta in nature, have outperformed more aggressive strategies. This was already shown by van Vliet and De Koning in their research between 1929-2008, but this study analyzes a different timeframe. In this instance, low-volatility portfolios outperformed high-volatility portfolios from 1968-2015, and low-beta portfolios outperformed high-beta portfolios from 1968-2015. The paper also looks at the characteristics of these portfolios in more detail by analyzing log-likelihood ratios that a stock selected at random is of a given style, like size or value, relative to the unconditional probability of being that style. The findings were very telling. On average, low-volatility stocks

were 30 times larger than high-volatility stocks. As a result, though the high-volatility names made up half the total number of stocks, they contributed to just 9% of the total market capitalization when compared with low volatility stocks. Moreover, when looking at the average returns across the portfolios of various volatilities, there seemed to be a relatively flat, slightly increased return with increasing volatility. The high-volatility/high-beta portfolio had a significant negative alpha, while the low-volatility/low-beta portfolio had a significant positive alpha.

Many of the volatile stocks tend to be small, unprofitable growth stocks, which can help explain the relationship of the defensive strategy performance to size, value, and profitability. In fact, since 1968-2015, the portfolio of high-volatility stocks almost mirrors the performance of the portfolio of unprofitable, small growth stocks. Thus, the outperformance of defensive stocks from 1968-2015 have delivered significant alphas, and can be explained when accounting for size, value, and profitability. In fact, profitability itself is so significant, that the case can be made that this alpha may be due in large part due to excluding unprofitable, small growth stocks.

1.2.2 Low-Risk Anomaly Applications

In this section, applications utilizing the Low-Risk Anomaly will be discussed. One example includes betting against Beta (BAB), which is a strategy used by quantitative hedge funds like AQR. Betting against Correlation (BAC) is a similar strategy in that it decomposes the effects of beta into two separate factors for a more concentrated investment. Moreover, index additions and deletions have proven to be instrumental in influencing the price of affected stocks, indicating that if a model can predict these movements accurately beforehand, there is an additional considerable arbitrage opportunity.

Betting Against Beta, by Andrea Frazzini and Lasse Heje Pedersen

One of the observations from prior literature looking at the Low-Risk Anomaly suggests that Beta is not a great measure of risk, and is more related to the anomaly than volatility is. AQR, a successful quantitative-focused hedge fund, employs a strategy called Betting Against Beta (BAB), which is a simple method of statistical arbitrage generated by shorting high-beta stocks and longing low-beta stocks (<http://www.investopedia.com/articles/investing/082515/how-aqr-places-bets-against-beta.asp>). As discussed in some of the works previously, the premise behind this arbitrage opportunity is that high-beta stocks are overpriced while low-beta stocks are underpriced. The theory is that the stocks will eventually return to this equilibrium point, called the security market lined (SML). While the prices approach this median, the investor can capture this spread as an arbitrage opportunity.

The Capital Asset Pricing Model (CAPM) calculates the expected equity return given certain levels of risk. Any excess above this risk-adjusted return is the Sharpe Ratio, or alpha that is generated by the stock. Investors try to maximize this number, and one way to do it is by leveraging up, or using borrowed capital to invest. By paying

a fixed interest rate on the borrowed capital, assuming the return is more than the interest rate, an investor can increase the return on their invested equity. As one can imagine, this is more risky, as returns in both the upward and downward direction are magnified. As a result, many investment managers like mutual funds are legally constrained on the amount of leverage they can apply to a portfolio. Due to this, many fund managers must overweight high-beta stocks to improve overall returns. This leads to a tilting towards beta, and a flattening of this SML in relation to CAPM. This leads to a pricing anomaly which firms like AQR can take advantage of. By creating a market neutral strategy, and shorting high-beta stocks while longing low-beta stocks, they can capture this opportunity.

In this paper, a real-world resembling model is created with leverage and margin constraints in 55,600 stocks from 20 global stock, bond, credit, and futures markets. Some agents in this model cannot use any leverage, and some have limited margin constraints, much like many investors and fund managers. As mentioned, many mutual funds, pension funds, and individual investors are constrained by the amount of leverage they can take on, such that instead of investing in a portfolio yielding the highest Sharpe Ratio, they are forced to overweight portfolios with higher-risk stocks. This suggests fund managers hold high-beta stocks to a lower risk adjusted return standard than low-beta stocks, which would require leverage. Thus, if one cannot leverage or has significant leverage constraints, then this agent will overweight riskier securities. The model in this paper was able to empirically show this in the equities, bonds, and futures markets. This was done by sorting portfolios by betas, and realizing alphas and Sharpe Ratios declining with increases in portfolio-beta.

Presumably, if one could lever up without constraint, the investor would underweight high-beta assets and overweight low-beta assets. BAB factors help explain this better. A BAB factor is a portfolio longing low-beta securities (leveraged to a beta of 1), shorting high-beta assets (deleveraged to a beta of 1), and is market neutral. The model in the paper predicts that this portfolio will have a positive return, that increases with the spread in the betas and tightness of leverage constraints. Thus, longing low-beta and shorting high-beta yields significant, and positive risk-adjusted returns. This was observed in the model by looking at U.S., developing, and international equity markets and observing that the BAB factor yielded a Sharpe ratio that was double its value effect, and 40% greater than momentum. The BAB factor had very high risk-adjusted returns, and during four twenty-year periods between 1926 and 2012, produced significant positive returns. This generally held across other asset classes, including credit and treasury bond markets.

When a leverage constraint was met or surpassed, and the agent needs to deleverage, the BAB factor portfolio experiences negative returns, but its expected future returns still increased. This was once again shown with a time series with spreads of various funding constraints. Another central idea of the model was that increased funding liquidity risk compresses betas toward one. This was proven by looking at the volatility of funding constraints as funding liquidity risk; the end result was that the dispersion of betas when funding liquidity risk is high, and was much lower than when funding liquidity risk is low. In other words, tightening of funding constraints leads to a lower BAB factor.

Finally, the model showed that investors that are more constrained are forced to overweight riskier securities, while investors without such constraints can overweight low-risk securities. Studying a number of stock portfolios from constrained investors, most fund managers and individual investors' portfolios have a beta greater than one. However, many private equity firms that perform a leveraged buyouts are traditionally able to purchase firms with a beta below one, and apply leverage, allowing them to utilize this anomaly since they have lower leverage constraints than their public market counterparts. Great investor Warren Buffett even bets against beta, as many of his investments are leveraged, low-beta stocks. By having these constraints, though, the typical investor is forced to hold on to riskier, high-beta stocks, leading to effective of the BAB factor.

Betting Against Correlation: Testing Theories of the Low-Risk Effect, by Cliff Asness, Andrea Frazzini, Niels Joachim Gormsen, and Lasse Heje Pedersen

As mentioned previously, the “low-risk effect” is the idea that lower-risk or lower-volatility stocks, tend to generate a higher alpha than higher-risk or higher-volatility stocks. In trying to understand the reason this anomaly occurs, Asness et al. consider two possible explanations. The first looks at whether this is caused by leverage constraints, or measurement using systematic risk. The second focuses on the behavioral effects, or idiosyncratic risks. One of the main issues with prior research, is that many low-risk factors are correlated and interrelated, making it hard to completely isolate certain factors or effects. In this paper, global data was used, with a couple new factors meant to control for existing factors.

“Financial Intermediaries and the Cross-Section of Asset Returns” by Adrian, Etula, and Muir in 2014, showed a link between return to the Betting Against Beta (BAB) factor and financial intermediary leverage (<http://faculty.som.yale.edu/tylermuir/documents/FINANCIALINTERMEDIARIESANDCROSSSECTIONASSETRETURNS.2013.pdf>). Many of these factors, including BAB, generally exhibited the “low-risk effect” and are were consequently very difficult to completely distinguish from one another. Thus, Asness, Frazzini, Gormsen, and Pedersen decided to attempt just that, by breaking down BAB into two separate factors: betting against correlation (BAC) and betting against volatility (BAV). This is done because beta itself can be decomposed into the stock's correlation with the market times its own volatility, divided by the market's overall volatility. BAC is accomplished through longing stocks with a low correlation to the market, and shorting those with a high correlation to the market, while trying to match the volatilities of both the long and short portfolios. BAV is achieved in a similar manner, except instead of longing and shorting correlation, volatility is longed and shorted while correlation is kept constant.

To address the behavioral explanation, the paper looks at some prior factors from observations made by Ang, Hodrick, Xing, and Zhang in their paper “The Cross-Section of Volatility and Expected Returns” (<http://onlinelibrary.wiley.com/doi/10.1111/j.1540-6261.2006.00836.x/abstract>). They found stocks with low idiosyncratic volatility (IVOL) had a greater risk-adjusted return, and in 2009 found

that a low maximum return (LMAX), a measure of idiosyncratic skewness, is associated with greater risk-adjusted returns (<http://web.ics.purdue.edu/~zhang654/jfe2009.pdf>). This paper being discussed kept the focus on the LMAX and IVOL, but added another factor, scaled MAX (SMAX), which longs stocks with a low MAX return divided by ex-ante volatility, and consequently shorts stocks with a high MAX return divided by ex-ante volatility. This allows focus on the lottery demand, holding volatility relatively constant and only focusing on the distribution of the returns. Margin debt held by investors, and investor sentiment were also noted.

Overall, 58,415 stocks from the MSCI World Index from 24 different countries between January 1926 and December 2015 were covered. BAV and BAC ended up being very successful in controlling for the other factors that could influence the “low-risk effect.” The BAV findings are in line with prior findings, and can be explained through behavioral effects like the lottery preference and leverage aversion. For all stocks, the BAC factor produced a significant six-factor alpha that was nearly independent of the other low-risk factors studied. This was partially due to the leverage aversion which indicates correlation changes in beta should be priced in. In terms of explaining the behavioral side with factors, SMAX was the only truly great, resilient measure used. The rest generally had higher turnover, and were consequently very susceptible to microstructure noise. SMAX attained positive risk-adjusted returns in the U.S. but negative risk-adjusted returns globally, which was seen with some other idiosyncratic risk factors. The paper showed that systematic low-risk factor generally tended to outperform behavioral risk factors, especially when considering turnover and time period length. All in all, the low-risk effect was believed to be driven by multiple factor effects, meaning both leverage constraints and the demand for lottery could play a role in influencing this. However, leverage constraint effects were a bit stronger, especially internationally. By cleaning up these factors, this paper is able to provide clearer explanations for previously observed results.

Price Response to Factor Index Additions and Deletions, by Joop Huij and Georgi Kyosev

One of the driving fundamental assumptions of finance is a flat demand curve for stocks, where risk is the main driver and each stock has a perfect substitute. However, this concept has been questioned for the past few years, with literature picking up on stocks showing supply shocks and quantifying how this affects their market price. Literature has shown several instances where large block sales of stock has negatively affected its price. This was often due to information contamination, which is new, significant information about the company in the market. This information often reflects fundamental changes in the company, and if it is negative, will understandably trigger block sales. Thus, the price change is less due to the supply shock, and more so due to the fundamental change in the company’s value (like a scandal or earnings report).

However, interesting patterns, that have not yet been fully explained, emerge from observations regarding S&P 500 company addition and deletions. When companies are added or removed from the index, it is purely mechanical, and

usually not due to some drastic fundamental change in the company. Assuming the market is efficient, the demand for stocks should not change due to being added or removed from an index, but several studies have shown that it does. Harris and Gurel (1986) (<http://onlinelibrary.wiley.com/doi/10.1111/j.1540-6261.1986.tb04550.x/full>), Shleifer (1986) (<http://onlinelibrary.wiley.com/doi/10.1111/j.1540-6261.1986.tb04518.x/abstract>), Beneish and Whaley (1996) (<http://onlinelibrary.wiley.com/doi/10.1111/j.1540-6261.1996.tb05231.x/full>), Chen, Noronha, and Singal (2004) (<http://onlinelibrary.wiley.com/doi/10.1111/j.1540-6261.2004.00683.x/full>) all show how new additions to the S&P come with higher than normal returns for that company. Though they agree on the price movement, the studies tend to disagree on the precise reason for this price movement; some possible explanations include compensation for providing liquidity, better monitoring for investors when a company is added to a reputable and large index, and higher analyst coverage leading to more information and analysis available on the company. One primary concern is whether or not index reshuffling is an information-free event - that is, whether a company being added or removed adds information to the market about the company.

In this paper, the authors look at factor index rebalancing for an information-free event. Factor indices are part of a parent index of many other stocks, and are constructed in a mechanical way that is publicly available and usually based on ranking stocks off a particular ratio of characteristic. Looking at the MSCI Minimum Volatility (USMV) index, returns were recorded for the stocks that had been added/dropped. It was found that the cumulative return from announcement to the effective day was 1.07% for stocks added with a significant t-statistic of 7.16, with 62% of the stocks exhibiting a positive cumulative abnormal return. Of the 1.07% increase, 0.63% of it was gained the following day, indicating that a large part of the increase is due to greater demand from index funds. 0.31% of the return is lost five days after the rebalancing, but generally the price tends to stabilize afterwards after ten days. Thus, 68% of the price increase is permanent, while the other 32% is temporary and lost after a few days. This can be due to a number of reasons including a liquidity premium charged by the stock's owner or arbitrage activity. Average trading volume was also significantly more for stocks that were recently added to the index. For the days between the announcement and actual addition of the stock, the average trading volume was 30% higher than normal, with a significant t-statistic of 3.81. Moreover, there is a 74% increase in volume for the day prior to the actual adding of the security. A very similar phenomena occurs with stocks set to be dropped from the USMV Index; from the announcement of a stock being dropped to the day before it is actually deleted from the index, the total cumulative abnormal return is -0.91%, and a total of -0.57% comes the day right before. After the stocks are deleted, 64% have a negative return the following day, and only 0.49% of the -0.91% is regained after three weeks. Trading volume also spikes 46% on the day prior to removal from the index. After three weeks, it returns back to within 1% of the normal trading volume.

These findings imply that once a security is added to a factor index, the demand curve shifts to the right, moving the equilibrium. The trading volume change is likely due to index funds buying or selling massive amounts of the stocks that will be added

or removed. Moreover, it was found that the amount of the return is also directly related to the weighting of the volume of stocks entering or leaving the factor index. All in all, these findings suggest an index arbitrage opportunity if the index additions or deletions can be predicted.

Chapter 2

Mathematics and Science

2.1 Math

T_EX is the best way to typeset mathematics. Donald Knuth designed T_EX when he got frustrated at how long it was taking the typesetters to finish his book, which contained a lot of mathematics. One nice feature of *R Markdown* is its ability to read LaTeX code directly.

If you are doing a thesis that will involve lots of math, you will want to read the following section which has been commented out. If you're not going to use math, skip over or delete this next commented section.

2.2 Chemistry 101: Symbols

Chemical formulas will look best if they are not italicized. Get around math mode's automatic italicizing in LaTeX by using the argument `$\mathrm{formula here}$` , with your formula inside the curly brackets. (Notice the use of the backticks here which enclose text that acts as code.)

So, Fe₂²⁺Cr₂O₄ is written `$\mathrm{Fe_2^{2+}Cr_2O_4}$` .

Exponent or Superscript: O⁻

Subscript: CH₄

To stack numbers or letters as in Fe₂²⁺, the subscript is defined first, and then the superscript is defined.

Bullet: CuCl • 7H₂O

Delta: Δ


Reaction Arrows: \longrightarrow or $\xrightarrow{\text{solution}}$

Resonance Arrows: \longleftrightarrow

Reversible Reaction Arrows: \rightleftharpoons


2.2.1 Typesetting reactions

You may wish to put your reaction in an equation environment, which means that LaTeX will place the reaction where it fits and will number the equations for you.



We can reference this combustion of glucose reaction via Equation (2.1).

2.2.2 Other examples of reactions



2.3 Physics

Many of the symbols you will need can be found on the math page <http://web.reed.edu/cis/help/latex/math.html> and the Comprehensive LaTeX Symbol Guide (<http://mirror.utexas.edu/ctan/info/symbols/comprehensive/symbols-letter.pdf>).

2.4 Biology

You will probably find the resources at <http://www.lecb.ncifcrf.gov/~toms/latex.html> helpful, particularly the links to bst files for various journals. You may also be interested in TeXShade for nucleotide typesetting (<http://homepages.uni-tuebingen.de/beitz/txe.html>). Be sure to read the proceeding chapter on graphics and tables.

Chapter 3

Tables, Graphics, References, and Labels

3.1 Tables

In addition to the tables that can be automatically generated from a data frame in **R** that you saw in [R Markdown Basics] using the `kable` function, you can also create tables using *pandoc*. (More information is available at <http://pandoc.org/README.html#tables>.) This might be useful if you don't have values specifically stored in **R**, but you'd like to display them in table form. Below is an example. Pay careful attention to the alignment in the table and hyphens to create the rows and columns.

Table 3.1: Correlation of Inheritance Factors for Parents and Child

| Factors | Correlation between Parents & Child | Inherited |
|-----------------------|-------------------------------------|-----------|
| Education | -0.49 | Yes |
| Socio-Economic Status | 0.28 | Slight |
| Income | 0.08 | No |
| Family Size | 0.18 | Slight |
| Occupational Prestige | 0.21 | Slight |

We can also create a link to the table by doing the following: Table 3.1. If you go back to [Loading and exploring data] and look at the `kable` table, we can create a reference to this max delays table too: Table ???. The addition of the `(\#tab:inher)` option to the end of the table caption allows us to then make a reference to Table `\@ref(tab:label)`. Note that this reference could appear anywhere throughout the document after the table has appeared.

3.2 Figures

If your thesis has a lot of figures, *R Markdown* might behave better for you than that other word processor. One perk is that it will automatically number the figures accordingly in each chapter. You'll also be able to create a label for each figure, add a caption, and then reference the figure in a way similar to what we saw with tables earlier. If you label your figures, you can move the figures around and *R Markdown* will automatically adjust the numbering for you. No need for you to remember! So that you don't have to get too far into LaTeX to do this, a couple **R** functions have been created for you to assist. You'll see their use below.

In the **R** chunk below, we will load in a picture stored as `reed.jpg` in our main directory. We then give it the caption of "Reed logo", the label of "reedlogo", and specify that this is a figure. Make note of the different **R** chunk options that are given in the R Markdown file (not shown in the knitted document).

```
include_graphics(path = "figure/reed.jpg")
```



Figure 3.1: Reed logo

Here is a reference to the Reed logo: Figure 3.1. Note the use of the `fig:` code here. By naming the **R** chunk that contains the figure, we can then reference that figure later as done in the first sentence here. We can also specify the caption for the figure via the R chunk option `fig.cap`.

Below we will investigate how to save the output of an **R** plot and label it in a way similar to that done above. Recall the `flights` dataset from Chapter ?? (Note that we've shown a different way to reference a section or chapter here.) We will next explore a bar graph with the mean flight departure delays by airline from Portland for 2014. Note also the use of the `scale` parameter which is discussed on the next page.

```
flights %>% group_by(carrier) %>%  
  summarize(mean_dep_delay = mean(dep_delay)) %>%  
  ggplot(aes(x = carrier, y = mean_dep_delay)) +  
  geom_bar(position = "identity", stat = "identity", fill = "red")
```




Figure 3.2: Mean Delays by Airline

Here is a reference to this image: Figure 3.2.

A table linking these carrier codes to airline names is available at <https://github.com/ismayc/pnwflights14/blob/master/data/airlines.csv>.

Next, we will explore the use of the `out.extra` chunk option, which can be used to shrink or expand an image loaded from a file by specifying `"scale= "`. Here we use the mathematical graph stored in the “subdivision.pdf” file.




Figure 3.3: Subdiv. graph

Here is a reference to this image: Figure 3.3. Note that `echo=FALSE` is specified so that the **R** code is hidden in the document.

More Figure Stuff

Lastly, we will explore how to rotate and enlarge figures using the `out.extra` chunk option. (Currently this only works in the PDF version of the book.)




Figure 3.4: A Larger Figure, Flipped Upside Down

As another example, here is a reference: Figure 3.4.

3.3 Footnotes and Endnotes

You might want to footnote something.¹ The footnote will be in a smaller font and placed appropriately. Endnotes work in much the same way. More information can be found about both on the CUS site or feel free to reach out to `data@reed.edu`.

3.4 Bibliographies

Of course you will need to cite things, and you will probably accumulate an armful of sources. There are a variety of tools available for creating a bibliography database (stored with the `.bib` extension). In addition to BibTeX suggested below, you may want to consider using the free and easy-to-use tool called Zotero. The Reed librarians have created Zotero documentation at <http://libguides.reed.edu/>

¹footnote text

citation/zotero. In addition, a tutorial is available from Middlebury College at <http://sites.middlebury.edu/zoteromiddlebury/>.

R Markdown uses *pandoc* (<http://pandoc.org/>) to build its bibliographies. One nice caveat of this is that you won't have to do a second compile to load in references as standard LaTeX requires. To cite references in your thesis (after creating your bibliography database), place the reference name inside square brackets and precede it by the “at” symbol. For example, here's a reference to a book about worrying: (Molina & Borkovec, 1994). This `Molina1994` entry appears in a file called `thesis.bib` in the `bib` folder. This bibliography database file was created by a program called BibTeX. You can call this file something else if you like (look at the YAML header in the main `.Rmd` file) and, by default, is to placed in the `bib` folder.

For more information about BibTeX and bibliographies, see our CUS site (<http://web.reed.edu/cis/help/latex/index.html>)². There are three pages on this topic: *bibtex* (which talks about using BibTeX, at <http://web.reed.edu/cis/help/latex/bibtex.html>), *bibtexstyles* (about how to find and use the bibliography style that best suits your needs, at <http://web.reed.edu/cis/help/latex/bibtexstyles.html>) and *bibman* (which covers how to make and maintain a bibliography by hand, without BibTeX, at <http://web.reed.edu/cis/help/latex/bibman.html>). The last page will not be useful unless you have only a few sources.

If you look at the YAML header at the top of the main `.Rmd` file you can see that we can specify the style of the bibliography by referencing the appropriate csl file. You can download a variety of different style files at <https://www.zotero.org/styles>. Make sure to download the file into the `csl` folder.

Tips for Bibliographies

- Like with thesis formatting, the sooner you start compiling your bibliography for something as large as thesis, the better. Typing in source after source is mind-numbing enough; do you really want to do it for hours on end in late April? Think of it as procrastination.
- The cite key (a citation's label) needs to be unique from the other entries.
- When you have more than one author or editor, you need to separate each author's name by the word “and” e.g. `Author = {Noble, Sam and Youngberg, Jessica},.`
- Bibliographies made using BibTeX (whether manually or using a manager) accept LaTeX markup, so you can italicize and add symbols as necessary.
- To force capitalization in an article title or where all lowercase is generally used, bracket the capital letter in curly braces.
- You can add a Reed Thesis citation³ option. The best way to do this is to use the `phdthesis` type of citation, and use the optional “type” field to enter “Reed thesis” or “Undergraduate thesis.”

²Reed College (2007)

³Noble (2002)

3.5 Anything else?

If you'd like to see examples of other things in this template, please contact the Data @ Reed team (email data@reed.edu) with your suggestions. We love to see people using *R Markdown* for their theses, and are happy to help.

Conclusion

If we don't want Conclusion to have a chapter number next to it, we can add the `{-}` attribute.

More info

And here's some other random info: the first paragraph after a chapter title or section head *shouldn't be* indented, because indents are to tell the reader that you're starting a new paragraph. Since that's obvious after a chapter or section title, proper typesetting doesn't add an indent there.

Appendix A

The First Appendix

This first appendix includes all of the R chunks of code that were hidden throughout the document (using the `include = FALSE` chunk tag) to help with readability and/or setup.

In the main Rmd file

In Chapter 3:

```
# This chunk ensures that the thesishdown package is  
# installed and loaded. This thesishdown package includes  
# the template files for the thesis and also two functions  
# used for labeling and referencing  
if(!require(devtools))  
  install.packages("devtools", repos = "http://cran.rstudio.com")  
if(!require(dplyr))  
  install.packages("dplyr", repos = "http://cran.rstudio.com")  
if(!require(ggplot2))  
  install.packages("ggplot2", repos = "http://cran.rstudio.com")  
if(!require(ggplot2))  
  install.packages("bookdown", repos = "http://cran.rstudio.com")  
if(!require(thesishdown)){  
  library(devtools)  
  devtools::install_github("ismayc/thesishdown")  
}  
library(thesishdown)  
flights <- read.csv("data/flights.csv")
```


Appendix B

The Second Appendix, for Fun

References

- Amadeo, K. (2017, april). Stock market crash of 2008. *The Balance*. Retrieved from <https://www.thebalance.com/stock-market-crash-of-2008-3305535>
- Angel, E. (2000). *Interactive computer graphics : A top-down approach with opengl*. Boston, MA: Addison Wesley Longman.
- Angel, E. (2001a). *Batch-file computer graphics : A bottom-up approach with quicktime*. Boston, MA: Wesley Addison Longman.
- Angel, E. (2001b). *Test second book by angel*. Boston, MA: Wesley Addison Longman.
- Baker, M., Bradley, B., & Wurgler, J. (2011). Benchmarks as limits to arbitrage: Understanding the low-volatility anomaly. *Financial Analysts Journal*, 67(1), 40–54.
- Berger, R. (2017, May). Amazon’s 49,000. *Forbes*. Forbes Magazine. Retrieved from <https://www.forbes.com/sites/robertberger/2017/05/20/amazons-49000-return-a-test-for-value-investors/#7b8efe2049cf>
- BlackRock. (2017, September). IShares edge msci min vol usa etf | usmv. Retrieved from <https://www.ishares.com/us/products/239695/ishares-msci-usa-minimum-volatility-etf>
- Boyer, B., Mitton, T., & Vorkink, K. (2009). Expected idiosyncratic skewness. *The Review of Financial Studies*, 23(1), 169–202.
- Brennan, M. J. (1989). Capital asset pricing model. In *Finance* (pp. 91–102). Springer.
- Fama, E. F., & French, K. R. (1992). The cross-section of expected stock returns. *The Journal of Finance*, 47(2), 427–465.
- Fischhoff, B., Slovic, P., & Lichtenstein, S. (1977). Knowing with certainty: The appropriateness of extreme confidence. *Journal of Experimental Psychology: Human Perception and Performance*, 3(4), 552.
- Fontinelle, E. (2009, October). ETF tracking errors: Protect your returns. Retrieved from <http://www.investopedia.com/articles/exchangetradedfunds/>

- 09/tracking-error-etf-funds.asp
- Goodwin, T. H. (1998). The information ratio. *Financial Analysts Journal*, 54(4), 34–43.
- Graham, B. (1965). *The intelligent investor: A book of practical counsel*. Prabhat Prakashan.
- Haugen, R. A., & Heins, J. A. (1975). Risk and the rate of return on financial assets: Some old wine in new bottles. *Journal of Financial and Quantitative Analysis*, 10(5), 775–784.
- Hayes, A. (2017, October). Exchange-traded fund (etf). Retrieved from <http://www.investopedia.com/terms/e/etf.asp>
- History. (2010). Stock market crash of 1929. *History.com*. A&E Television Networks. Retrieved from <http://www.history.com/topics/1929-stock-market-crash>
- Investopedia. (2015, April). What are some common measures of risk used in risk management? *Investopedia*. Retrieved from <http://www.investopedia.com/ask/answers/041415/what-are-some-common-measures-risk-used-risk-management.asp>
- Investopedia. (2017, February). Two and twenty. *Investopedia*. Retrieved from http://www.investopedia.com/terms/t/two_and_twenty.asp
- Jensen, M. C., Black, F., & Scholes, M. S. (1972). The capital asset pricing model: Some empirical tests.
- Karceski, J. (2002). Returns-chasing behavior, mutual funds, and beta's death. *Journal of Financial and Quantitative Analysis*, 37(4), 559–594.
- MacDonald, L. (2009, october). How traders are front-running etfs. Retrieved from <https://seekingalpha.com/article/165877-how-traders-are-front-running-etfs>
- Molina, S. T., & Borkovec, T. D. (1994). The Penn State worry questionnaire: Psychometric properties and associated characteristics. In G. C. L. Davey & F. Tallis (Eds.), *Worrying: Perspectives on theory, assessment and treatment* (pp. 265–283). New York: Wiley.
- MSCI. (2013, September). Barra open optimizer. Retrieved from https://www.msci.com/documents/10199/242721/Barra_Optimizer.pdf/5acf4515-dbd8-4d0c-8acc-ca70b7b68f69
- Mullins, D. W. (1982). Does the capital asset pricing model work. *Harvard Business Review*, 60(1), 105–113.
- Noble, S. G. (2002). *Turning images into simple line-art* (Undergraduate thesis).

Reed College.

Ontario, S. C. (2017, May). Types of investment risk | understanding risk. Retrieved from <https://www.getsmarteraboutmoney.ca/invest/investing-basics/understanding-risk/types-of-investment-risk/>

Pettengill, G. N., Sundaram, S., & Mathur, I. (1995). The conditional relation between beta and returns. *Journal of Financial and Quantitative Analysis*, 30(1), 101–116.

Reed College. (2007, march). LaTeX your document. Retrieved from <http://web.reed.edu/cis/help/LaTeX/index.html>

Sensoy, B. A. (2009). Performance evaluation and self-designated benchmark indexes in the mutual fund industry. *Journal of Financial Economics*, 92(1), 25–39.

Vliet, P. V., & Koning, J. de. (2016). *High returns from low risk: A remarkable stock market paradox*. John Wiley & Sons.

Wynne, B. (2016, april). Understanding the iShares msci usa minimum volatility etf. Retrieved from <https://seekingalpha.com/article/3964639-understanding-ishares-msci-usa-minimum-volatility-etf>